

bank-loan-1

December 24, 2025

1 BANK LOAN APPROVAL PREDICTION - LOGISTIC REGRESSION PROJECT

Predict whether a bank should approve a customer's loan

Goal:- Predict Approved=1 or Not Approved=0 using customer features

2 Project Covers:

Clean dataset loading

Feature selection

Train/Test split

Scaling

Training(Logistic Regression)

Evaluation (Accuracy, Precision, Recall, F1)

Confusion Matrix + Classification Report

Decision Thresholds

AUC-ROC Curve

Prediction function

3 Import libraries

```
[42]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import(
    accuracy_score, ## Accuracy: Out of all predictions, how many were correct?
```

```

precision_score,## Precision: Out of all YES predictions, how many were
↪actually YES?
recall_score,# Recall: Out of all actual YES cases, how many did the model
↪find?
f1_score,# F1 Score: A single score that balances Precision and Recall, # A
↪table that shows Correct and wrong predictions (YES/NO)
confusion_matrix, # A table that shows Correct and wrong predictions (YES/
↪NO)
classification_report, # Shows Accuracy, Precision, Recall, and F1 together
↪in one report
roc_curve, # A graph that shows how well the model separates YES and NO
roc_auc_score# A number that tells how good the model is at separating
↪classes
)

```

4 LOAD DATASET

```
[43]: df=pd.read_csv('bank_loan.csv')
print(df.head())
print('\nShape',df.shape)
print('\nInfo',df.info())
print(df.isnull().sum())
```

	Age	Income	CreditScore	EmploymentStatus	Approved
0	59	40358	812	Unemployed	0
1	49	23267	595	Unemployed	0
2	35	102745	619	Salaried	1
3	63	109588	871	Self-employed	0
4	28	58513	648	Salaried	1

Shape (500, 5)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 5 columns):

#	Column	Non-Null Count	Dtype
0	Age	500 non-null	int64
1	Income	500 non-null	int64
2	CreditScore	500 non-null	int64
3	EmploymentStatus	500 non-null	object
4	Approved	500 non-null	int64

dtypes: int64(4), object(1)
memory usage: 19.7+ KB

Info None

Age 0

```
Income          0
CreditScore     0
EmploymentStatus 0
Approved        0
dtype: int64
```

```
[44]: #Count each employment status category
print('\n Employment Status Distribution')
print(df['EmploymentStatus'].value_counts())

#Previous last 10 rows
print('\n Last 10 rows of the dataset:')
print(df.tail(10))
```

```
Employment Status Distribution
EmploymentStatus
Salaried      289
Self-employed 136
Unemployed    75
Name: count, dtype: int64
```

```
Last 10 rows of the dataset:
   Age  Income  CreditScore EmploymentStatus  Approved
490  43     32219           870  Self-employed      0
491  37     20235           662  Self-employed      0
492  46     62929           856  Self-employed      0
493  28     147309          644  Salaried         1
494  49     87444            469  Salaried         1
495  46     159639          655  Salaried         1
496  30     121834          490  Salaried         0
497  46     104555          721  Salaried         1
498  54     93698            306  Unemployed       0
499  61     144450          432  Salaried         0
```

```
[45]: print(df.describe())

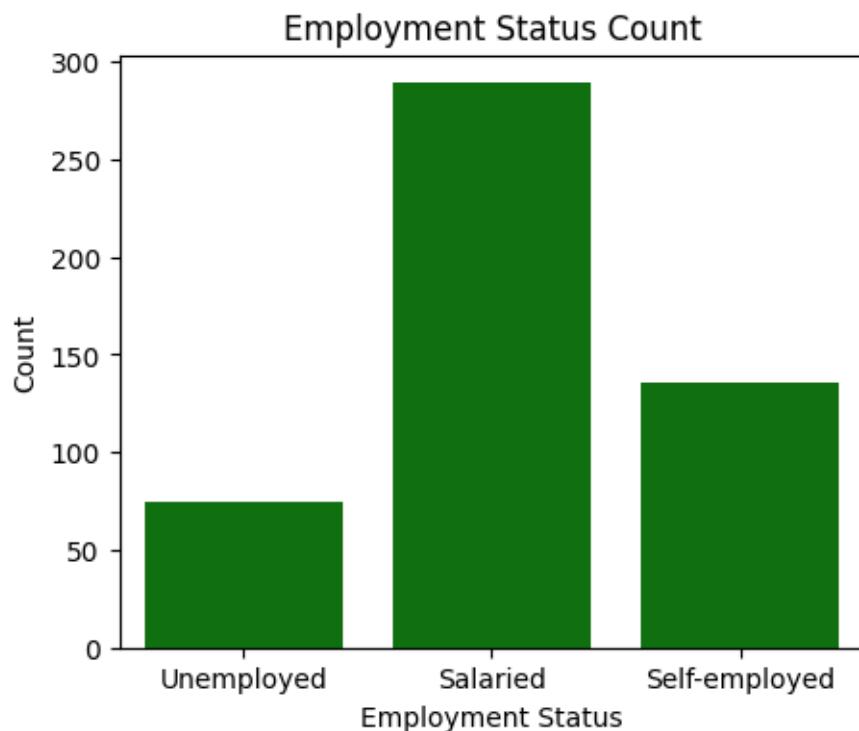
print('\n Unique Employment Status')
print(df['EmploymentStatus'].unique())
```

	Age	Income	CreditScore	Approved
count	500.000000	500.000000	500.000000	500.000000
mean	43.116000	109045.696000	583.490000	0.502000
std	12.733217	52974.205023	171.614805	0.500497
min	21.000000	20235.000000	300.000000	0.000000
25%	32.000000	62479.250000	440.000000	0.000000
50%	44.000000	109228.500000	585.500000	1.000000
75%	53.000000	156195.000000	723.500000	1.000000

```
max      64.000000 199208.000000  898.000000  1.000000
```

```
Unique Employment Status  
['Unemployed' 'Salaried' 'Self-employed']
```

```
[46]: #4.Employment status count Plat  
plt.figure(figsize=(5,4))  
sns.countplot(x='EmploymentStatus',data=df,color='Green')  
plt.title('Employment Status Count')  
plt.xlabel('Employment Status')  
plt.ylabel('Count')  
plt.show()
```



5 Feature Selection And Cleaning

```
[47]: df=df[['Age','Income','CreditScore','EmploymentStatus','Approved']]
```

6 Encode Employment Status(numeric encoding)

```
[48]: from sklearn.preprocessing import LabelEncoder  
le=LabelEncoder()  
df['EmploymentEncoded']=le.fit_transform(df['EmploymentStatus'])
```

```
[49]: #Final Feature/Target  
X=df[['Age','Income','CreditScore','EmploymentEncoded']]  
y=df['Approved']  
print('\n final features:\n',X.head())
```

```
final features:  
   Age  Income  CreditScore  EmploymentEncoded  
0    59     40358          812                  2  
1    49     23267          595                  2  
2    35     102745         619                  0  
3    63     109588         871                  1  
4    28     58513          648                  0
```

7 TRAIN-TEST SPLIT

```
[50]: X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.  
        ↪25,random_state=42,stratify=y)  
print('train size:',X_train.shape)  
print('test size:',X_test.shape)
```

```
train size: (375, 4)  
test size: (125, 4)
```

8 Feature Scaling

```
[51]: Scaler=StandardScaler()  
X_train=Scaler.fit_transform(X_train)  
X_test=Scaler.transform(X_test)
```

9 Train Logistic Regression Model

```
[52]: model=LogisticRegression(max_iter=1000)  
      ↪#Improves Accuracy.#Create a logistic  
      ↪regression model  
model.fit(X_train,y_train)  
      ↪#Train the model using scaled training data and their  
      ↪correct labels  
print('Model Training Complete')
```

```
Model Training Complete
```

10 Model Predictions

```
[53]: y_pred=model.predict(X_test) #Predicts class labels(0 or 1)
y_prob=model.predict_proba(X_test)[:,1] #Get probability of class 1(Loan
    ↪Approved) for each test sample
```

11 Evaluation Metrics

```
[54]: print("\n-----MODEL EVALUATION-----")
print("\nAccuracy Score:",accuracy_score(y_test,y_pred)) #How many total
    ↪predictions are correct
print("\nPrecision Score:",precision_score(y_test,y_pred)) #Out of all
    ↪predicted YES, how many were actually Yes(Controls False Positive)
print("\nRecall Score:",recall_score(y_test,y_pred)) #Out of all actual YES, how
    ↪many the model correctly found (Controls False Positive)
print("\nF1 Score:",f1_score(y_test,y_pred)) #Balance of precision and recall
    ↪in one score
```

-----MODEL EVALUATION-----

Accuracy Score: 0.624

Precision Score: 0.6212121212121212

Recall Score: 0.6507936507936508

F1 Score: 0.6356589147286822

12 Model Evaluation

Accuracy = 0.62 (62%) Out of 100 loan predictions, about 62 are correct.

Precision = 0.62 (62%) When the model says “YES - loan approved”, it is correct 62 times out of 100.

Recall = 0.65 (65%) Out of all people who should get the loan, the model correctly finds 65% of them.

F1 Score = 0.63(63%) Overall balanced performance between Precision and Recall.

13 Confusion Matrix + Classification Report

```
[55]: print("\nConfusion Matrix:")
print(confusion_matrix(y_test,y_pred))
#Shows TP,Fp,TN,FN to understand correct and incorrect predictions
```

Confusion Matrix:

```
[[37 25]
 [22 41]]
```

14 Confusion Matrix

37 → Correctly predicted NO

25 → Wrongly predicted YES (but actually NO)

22 → Wrongly predicted NO (but actually YES)

41 → Correctly predicted YES

```
[56]: print("\nClassification Report:")
print(classification_report(y_test,y_pred))
#Shows Precision,Recall,F1 score and support for each class
```

Classification Report:

	precision	recall	f1-score	support
0	0.63	0.60	0.61	62
1	0.62	0.65	0.64	63
accuracy			0.62	125
macro avg	0.62	0.62	0.62	125
weighted avg	0.62	0.62	0.62	125

15 Class 0 (NO / Loan Rejected)

Precision 63% → When model says NO, it is correct 63 times out of 100

Recall 60% → Found 60% of actual NO cases

F1-score 61% → Balanced score for NO class

Support 62 → Total 62 NO records

16 Class 1 (YES / Loan Approved)

Precision 62% → When model says YES, it is correct 62 times out of 100

Recall 65% → Found 65% of actual YES cases F1-score 64% → Balanced score for YES class

Support 63 → Total 63 YES records

Accuracy = 62% → Model predicted 62 correct out of 100

Macro Avg → Simple average of both classes

Weighted Avg → Average based on number of records

17 DECISION THRESHOLD ANALYSIS

Decision Threshold Analysis means deciding at what probability value the model should say YES or NO.

=>0.5 -> Yes

< 0.5 -> No

```
[57]: threshold=0.5;#Equal importance to YES and NO / Balanced datasets

print("\n Default Threshold=",threshold)

print("Prediction Probabilities(First 10):")

print(y_prob[:10]); #Shows the first 10 probability values predict by the model
```

```
Default Threshold= 0.5
Prediction Probabilities(First 10):
[0.32716288 0.68291553 0.52806798 0.675654 0.38418969 0.71675992
 0.69488846 0.63420475 0.59271349 0.5549581 ]
```

18 AUC-ROC Curve

AUC_ROC tells us how well a model can separate positive and negative cases.

It measures how good the model is at distinguishing YES vs NO.

```
[58]: # AUC tells how well the model separates Yes vs No.

# Step 1: Calculate False Positive Rate, True Positive Rate and Thresholds
fpr, tpr, thresholds = roc_curve(y_test,y_prob)
# fpr = how many wrong positives at each threshold
# tpr = how many correct positives at each threshold
# thresholds = different probability cut-off values

# Step 2: Calculate AUC score
auc = roc_auc_score(y_test,y_prob)
# AUC tells how well the model separates class 0 and class 1
# Higher AUC = better model
print(" \nAUC Score:", auc)
```

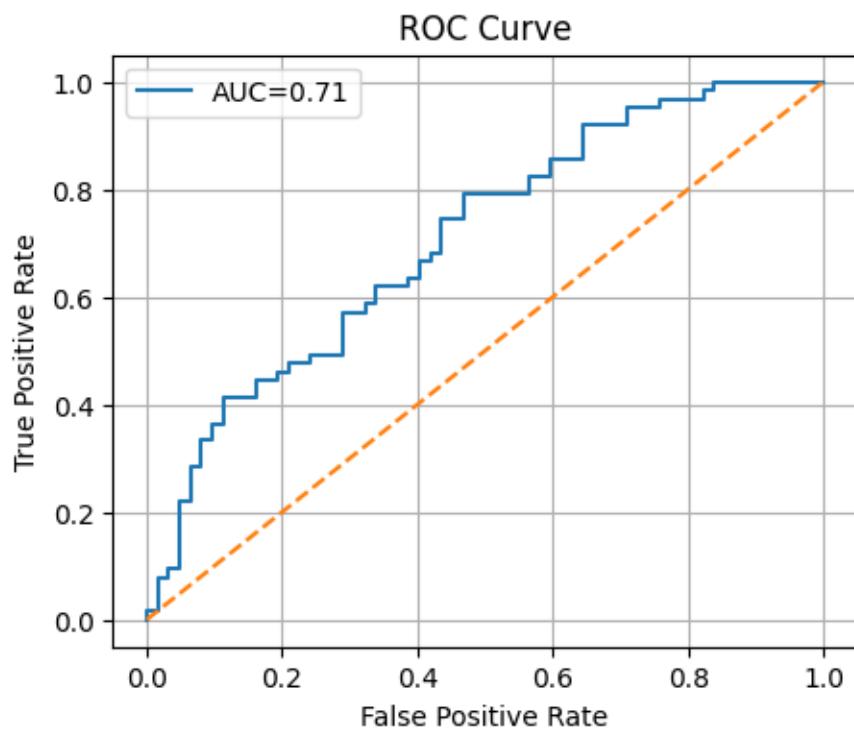
AUC Score: 0.7073732718894009

19 AUC Score = 0.71

Means 71% chance the model will correctly separates Yes vs No

```
[59]: plt.figure(figsize=(5,4))
plt.plot(fpr,tpr,label=f"AUC={auc:.2f}")
#Plot the ROC CURVE

plt.plot([0,1],[0,1], '--') #Draws a random guessing line / Used for comparison
plt.title('ROC Curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.grid(True)
plt.show()
```



20 PREDICTION FUNCTION

```
[60]: mapping = {'Salaried': 0, 'Self-employed': 1, 'Unemployed': 2}

def predict_loan(age,income,credit,employment):
    #1. Put the customer data into the DataFrame
```

```

row=pd.
>DataFrame([[age,income,credit,mapping[employment]]],columns=['Age','Income','CreditScore',''])

#2. Scale the data(same scaling used during training)
row_scaled=Scaler.transform(row);

#3. Get probability of loan approved
prob=model.predict_proba(row_scaled)[0][1]

#4.Final prediction (0 or 1) using threshold 0.5
pred=1 if prob >= 0.5 else 0;

#5. Print Simple Output
print(f"\n Loan Approval Probability={prob:.2f}")

if pred == 1:
    print('Loan Approved')
else:
    print('Loan Not Approved')

predict_loan(35,190000,800,'Salaried')
predict_loan(55,30000,600,'Unemployed')

```

Loan Approval Probability=0.88

Loan Approved

Loan Approval Probability=0.17

Loan Not Approved